

Programmatic Generation

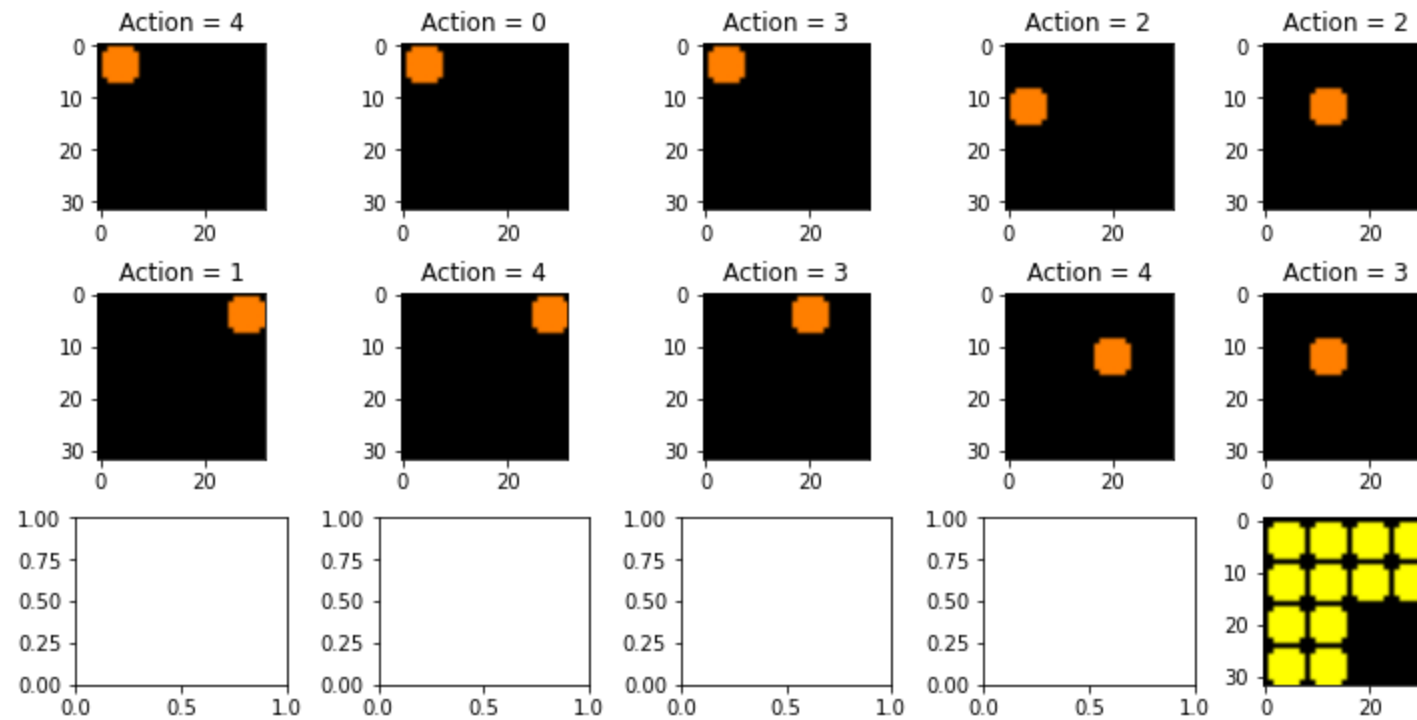
Recap

Slot attention + a simple MLP seems to be able to learn the dynamics of the shapeworld environment for certain instances.

This Week: Focused on where this fails, how this fails under the "single sphere" dataset.

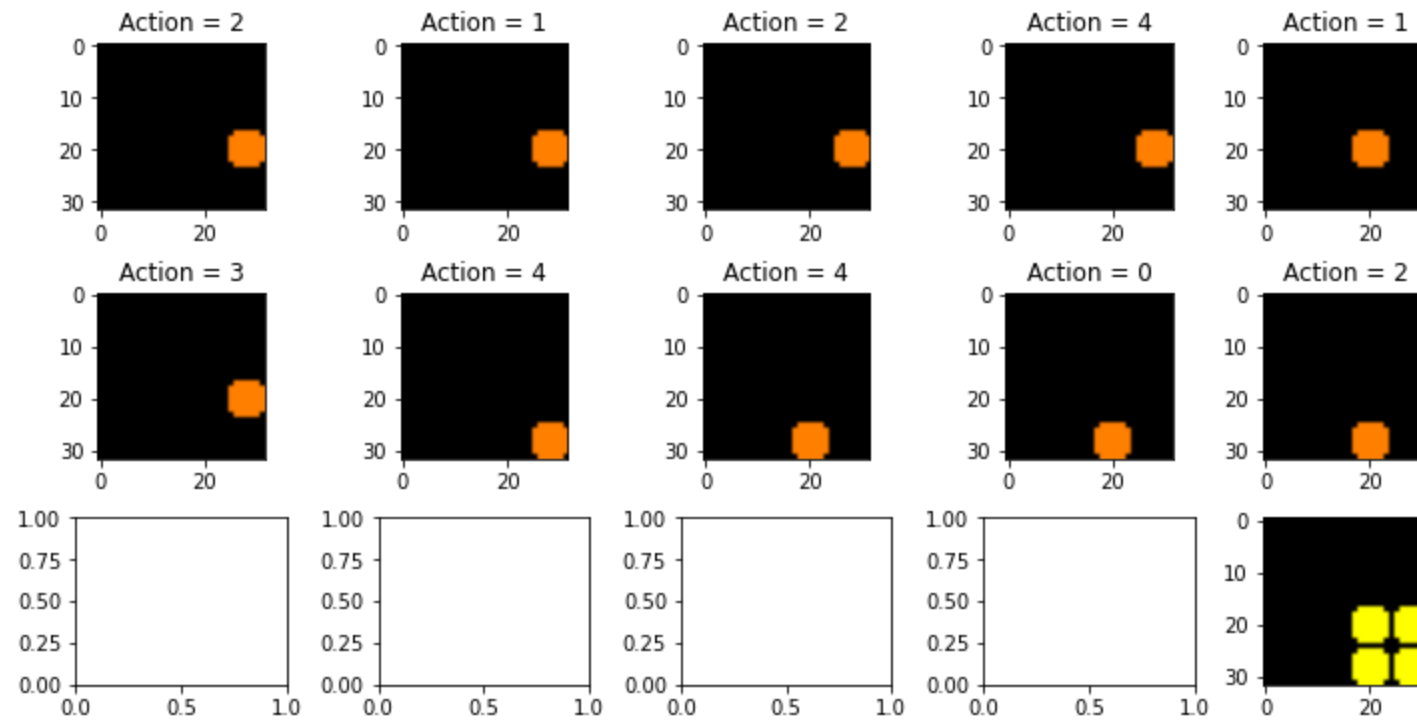
Single Sphere Dataset

Train time:



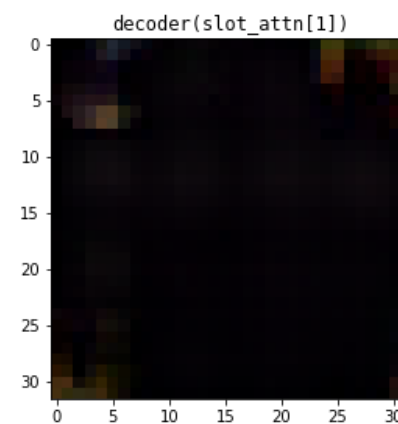
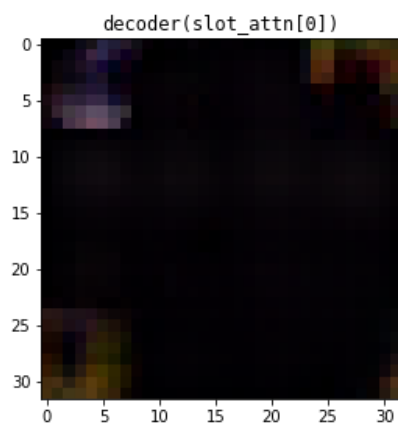
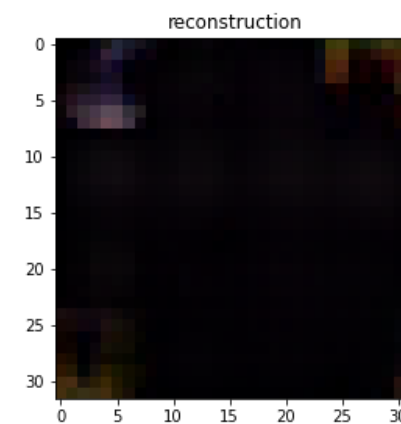
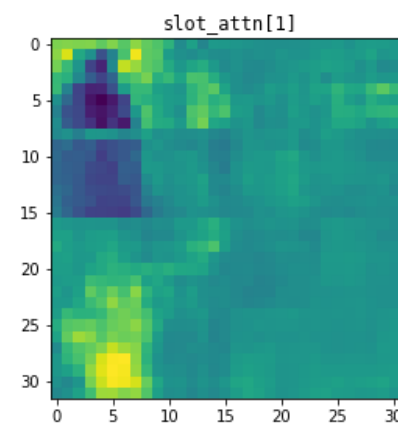
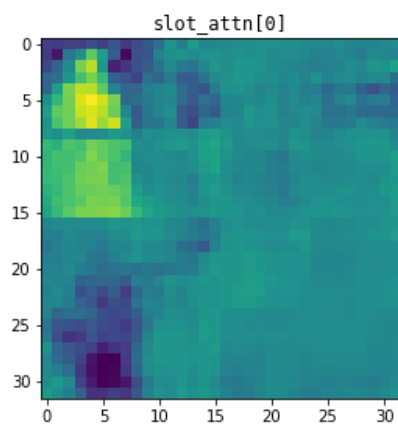
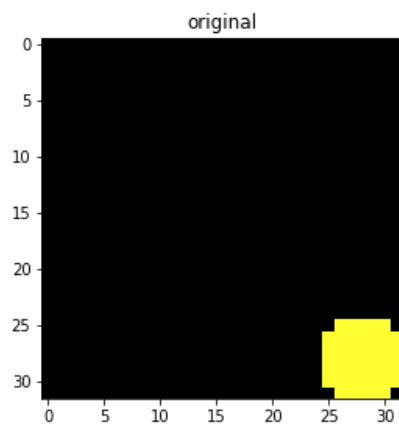
Single Sphere Dataset

Test time:



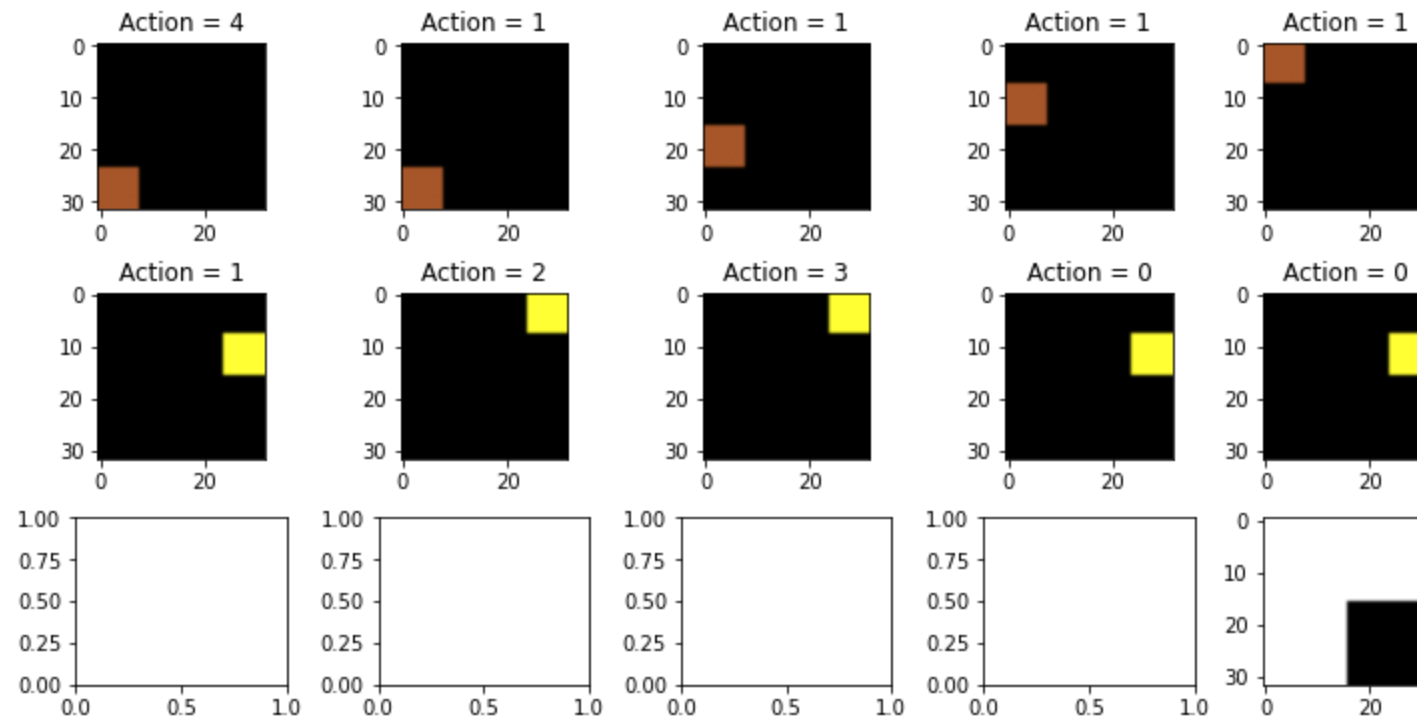
Overfitting

SlotVAEModule-6 (epoch=99, lr=3e-4, reg=1e-6)



Single Shapes Dataset

Train time:



Single Shapes Dataset

Test time:

